

Supplementary Material

From Laboratory to Field: A Leakage-Controlled Cross-Domain Evaluation Protocol for Plant Disease Recognition

Anonymous Author(s)

Affiliation withheld for double-blind review

This supplement carries the complete metric tables behind the compact versions in the main paper, the per-seed values every aggregate is computed from, and the per-run diagnostic figures. Nothing here is a new result: it is the same artifacts at full width.

All values come from `experiments/ica26/metrics/*.json` via `scripts/ica26_build_tables.py` and `scripts/ica26_sig`. The corpus is bound to dataset lock 83449f6013691855. This version covers 18 of 18 planned runs across seeds 42, 1337, 2026.

1 Dataset Provenance and Leakage Control

Table 1. Dataset provenance and leakage control. Perceptual-hash Hamming threshold 6; all 16 near pairs human-adjudicated as clearly different.

Quantity	Value
PlantDoc acquired (source)	2 578
PlantDoc previous effective	2 564
PlantDoc Core	2 561
PlantDoc non-reviewed unchanged	2 554
PlantDoc retained reviewed	7
PlantDoc excluded reviewed	17
Cross-dataset exact duplicates (acquired)	0
Cross-dataset near duplicates (acquired)	16
Cross-dataset exact duplicates (Core)	0
Cross-dataset near duplicates (Core)	16
Unresolved leakage pairs	0
Shared cross-domain classes	21
Cross-domain evaluation images	1 951

2 Complete Performance Tables

Table 2. In-domain test performance, mean $\pm$ sample standard deviation over independent training seeds (42, 1337, 2026). Macro averages are computed over the classes present in the evaluation split; PlantDoc Core’s arthropod-pest class has zero test images and is reported separately rather than averaged in as zero.

Model	Dataset	Seeds	N	Accuracy	Macro-F1	Weighted-F1	Balanced Acc.	Macro Prec.	Macro Rec.
ResNet-50	PlantVillage	3	10 709	0.9945 $\pm$ 0.0021	0.9927 $\pm$ 0.0029	0.9945 $\pm$ 0.0021	0.9923 $\pm$ 0.0029	0.9934 $\pm$ 0.0028	0.9923 $\pm$ 0.0029
EfficientNet-B0	PlantVillage	3	10 709	0.9964 $\pm$ 0.0005	0.9950 $\pm$ 0.0006	0.9964 $\pm$ 0.0005	0.9945 $\pm$ 0.0008	0.9956 $\pm$ 0.0004	0.9945 $\pm$ 0.0008
MobileNetV3-Small	PlantVillage	3	10 709	0.9948 $\pm$ 0.0004	0.9927 $\pm$ 0.0004	0.9948 $\pm$ 0.0004	0.9916 $\pm$ 0.0003	0.9940 $\pm$ 0.0006	0.9916 $\pm$ 0.0003
ResNet-50	PlantDoc Core	3	225	0.6637 $\pm$ 0.0403	0.6823 $\pm$ 0.0386	0.6946 $\pm$ 0.0406	0.6551 $\pm$ 0.0393	0.7539 $\pm$ 0.0317	0.6551 $\pm$ 0.0393
EfficientNet-B0	PlantDoc Core	3	225	0.6652 $\pm$ 0.0128	0.6768 $\pm$ 0.0139	0.6920 $\pm$ 0.0162	0.6569 $\pm$ 0.0087	0.7430 $\pm$ 0.0082	0.6569 $\pm$ 0.0087
MobileNetV3-Small	PlantDoc Core	3	225	0.5600 $\pm$ 0.0044	0.5876 $\pm$ 0.0061	0.5936 $\pm$ 0.0044	0.5640 $\pm$ 0.0030	0.6564 $\pm$ 0.0136	0.5640 $\pm$ 0.0030

Table 3. Cross-domain generalisation, mean $\pm$ sample standard deviation over three seeds: PlantVillage-trained models evaluated on the frozen shared-class PlantDoc Core subset. Metrics are computed only over the 21 shared classes; unmapped PlantDoc images are excluded from the evaluation set rather than counted as errors. Retained probability mass is the share of the source model’s belief falling inside the shared space before renormalisation.

Model	Dataset	Seeds	N	Accuracy	Macro-F1	Weighted-F1	Balanced Acc.	Macro Prec.	Macro Rec.	Shared classes	Retained prob. mass
ResNet-50	PlantVillage $\rightarrow$ PlantDoc Core	3	1 951	0.2583 ± 0.0206	0.2291 ± 0.0107	0.2512 ± 0.0138	0.2389 ± 0.0188	0.3629 ± 0.0197	0.2389 ± 0.0188	21	0.6367 ± 0.0945
EfficientNet-B0	PlantVillage $\rightarrow$ PlantDoc Core	3	1 951	0.2653 ± 0.0102	0.2300 ± 0.0065	0.2528 ± 0.0078	0.2475 ± 0.0110	0.3455 ± 0.0257	0.2475 ± 0.0110	21	0.6595 ± 0.0086
MobileNetV3-Small	PlantVillage $\rightarrow$ PlantDoc Core	3	1 951	0.2238 ± 0.0162	0.1869 ± 0.0109	0.2013 ± 0.0160	0.2141 ± 0.0135	0.3090 ± 0.0248	0.2141 ± 0.0135	21	0.5662 ± 0.0477

Table 4. Domain-shift degradation, mean $\pm$ sample standard deviation over three seeds. Drops are computed within each seed and then averaged, preserving the pairing between a checkpoint and its own cross-domain evaluation. In-domain is the PlantVillage test split over 38 classes; cross-domain is the PlantDoc Core shared-class subset over 21 classes. The two label spaces differ, so the drop combines domain shift with the change of task difficulty and should be read as a paired trend across models rather than as an isolated quantity.

Model	Seeds	Accuracy (in-domain)	Accuracy (cross-domain)	Accuracy abs. drop	Accuracy rel. drop %	Macro-F1 (in-domain)	Macro-F1 (cross-domain)	Macro-F1 abs. drop	Macro-F1 rel. drop %
ResNet-50	3	0.9945 $\pm$ 0.0021	0.2583 $\pm$ 0.0206	0.7361 $\pm$ 0.0213	74.02 $\pm$ 2.09	0.9927 $\pm$ 0.0029	0.2291 $\pm$ 0.0107	0.7636 $\pm$ 0.0131	76.92 $\pm$ 1.13
EfficientNet-B0	3	0.9964 $\pm$ 0.0005	0.2653 $\pm$ 0.0102	0.7310 $\pm$ 0.0097	73.37 $\pm$ 1.01	0.9950 $\pm$ 0.0006	0.2300 $\pm$ 0.0065	0.7650 $\pm$ 0.0064	76.89 $\pm$ 0.65
MobileNetV3-Small	3	0.9948 $\pm$ 0.0004	0.2238 $\pm$ 0.0162	0.7710 $\pm$ 0.0162	77.50 $\pm$ 1.63	0.9927 $\pm$ 0.0004	0.1869 $\pm$ 0.0109	0.8058 $\pm$ 0.0111	81.17 $\pm$ 1.10

Table 5. Model efficiency, mean $\pm$ sample standard deviation over three seeds. Epoch counts vary between seeds because early stopping fires at different points, so training time is reported alongside seconds per epoch. Latency and throughput are measured after warm-up on the accelerator reported in the run metadata.

Model	Dataset	Seeds	Params (M)	Epochs	Train time (min)	s / epoch	Latency b1 (ms)	Throughput (img/s)	Peak mem (MB)
ResNet-50	PlantVillage	3	23.59	12.0 $\pm$ 2.6	143.9 $\pm$ 35.1	718.7 $\pm$ 53.1	8.82 $\pm$ 0.36	216.5 $\pm$ 5.7	9472.5 $\pm$ 18.5
EfficientNet-B0	PlantVillage	3	4.06	13.0 $\pm$ 1.7	127.8 $\pm$ 16.1	590.4 $\pm$ 21.2	9.26 $\pm$ 0.22	271.7 $\pm$ 5.7	9241.6 $\pm$ 6.9
MobileNetV3-Small	PlantVillage	3	1.56	13.0 $\pm$ 3.5	30.4 $\pm$ 7.1	141.2 $\pm$ 7.1	6.42 $\pm$ 0.31	1286.8 $\pm$ 22.8	2150.0 $\pm$ 9.2
ResNet-50	PlantDoc Core	3	23.57	10.7 $\pm$ 2.1	11.3 $\pm$ 2.2	63.9 $\pm$ 7.6	8.58 $\pm$ 0.21	209.9 $\pm$ 5.7	11849.1 $\pm$ 591.2
EfficientNet-B0	PlantDoc Core	3	4.04	9.3 $\pm$ 0.6	10.6 $\pm$ 1.9	67.8 $\pm$ 11.4	9.75 $\pm$ 0.17	192.6 $\pm$ 131.2	11829.6 $\pm$ 339.5
MobileNetV3-Small	PlantDoc Core	3	1.55	13.0 $\pm$ 2.0	2.7 $\pm$ 0.4	12.6 $\pm$ 0.3	6.28 $\pm$ 0.23	1268.4 $\pm$ 5.1	2150.7 $\pm$ 5.3

3 Efficiency

4 Complete Calibration and Selective Prediction

5 Ranking Stability Across Seeds

6 Statistical Tests

7 Training Curves

Every seed is drawn rather than averaged. Early stopping fires at different epochs, so a mean curve over epoch index would silently truncate to the earliest stop and misrepresent the runs that trained longer.




Fig. 1. Training loss and validation macro-F1. Colour is the backbone, dash pattern is the seed.

Table 6. Confidence calibration and selective prediction, mean $\pm$ sample standard deviation over three seeds. Temperature is fitted on the held-out validation split only and applied unchanged to test and cross-domain logits; it is never refitted on evaluation data. AUC is the area under the risk-coverage curve; lower is better.

Model	Dataset	Setting	Seeds	LL	Brier	ECE	MCE	AUC	T
ResNet-50	PlantVillage	in-domain	3	0.1218 $\pm$ 0.0138	0.0195 $\pm$ 0.0042	0.0987 $\pm$ 0.0113	0.3504 $\pm$ 0.0611	0.0007 $\pm$ 0.0002	1.0000 $\pm$ 0.0000
ResNet-50	PlantVillage	in-domain (T-scaled)	3	0.0283 $\pm$ 0.0041	0.0092 $\pm$ 0.0026	0.0092 $\pm$ 0.0011	0.3061 $\pm$ 0.0897	0.0006 $\pm$ 0.0002	0.7028 $\pm$ 0.0076
ResNet-50	PlantVillage	cross-domain	3	2.9441 $\pm$ 0.1055	0.9938 $\pm$ 0.0342	0.2829 $\pm$ 0.0459	0.5042 $\pm$ 0.0136	0.6428 $\pm$ 0.0282	1.0000 $\pm$ 0.0000
ResNet-50	PlantVillage	cross-domain (T-scaled)	3	3.5879 $\pm$ 0.2283	1.1068 $\pm$ 0.0529	0.4178 $\pm$ 0.0533	0.6400 $\pm$ 0.0236	0.6427 $\pm$ 0.0282	0.7028 $\pm$ 0.0076
EfficientNet-B0	PlantVillage	in-domain	3	0.1195 $\pm$ 0.0075	0.0165 $\pm$ 0.0019	0.1015 $\pm$ 0.0069	0.5075 $\pm$ 0.1154	0.0009 $\pm$ 0.0003	1.0000 $\pm$ 0.0000
EfficientNet-B0	PlantVillage	in-domain (T-scaled)	3	0.0230 $\pm$ 0.0017	0.0059 $\pm$ 0.0006	0.0095 $\pm$ 0.0013	0.4969 $\pm$ 0.1793	0.0007 $\pm$ 0.0003	0.7006 $\pm$ 0.0040
EfficientNet-B0	PlantVillage	cross-domain	3	2.8680 $\pm$ 0.0086	0.9588 $\pm$ 0.0242	0.2561 $\pm$ 0.0186	0.4783 $\pm$ 0.0577	0.6156 $\pm$ 0.0280	1.0000 $\pm$ 0.0000
EfficientNet-B0	PlantVillage	cross-domain (T-scaled)	3	3.4624 $\pm$ 0.0278	1.0665 $\pm$ 0.0314	0.3973 $\pm$ 0.0192	0.5966 $\pm$ 0.0337	0.6152 $\pm$ 0.0295	0.7006 $\pm$ 0.0040
MobileNetV3-Small	PlantVillage	in-domain	3	0.1213 $\pm$ 0.0032	0.0199 $\pm$ 0.0015	0.0988 $\pm$ 0.0025	0.6900 $\pm$ 0.1577	0.0003 $\pm$ 0.0000	1.0000 $\pm$ 0.0000
MobileNetV3-Small	PlantVillage	in-domain (T-scaled)	3	0.0271 $\pm$ 0.0012	0.0086 $\pm$ 0.0006	0.0091 $\pm$ 0.0003	0.4771 $\pm$ 0.2112	0.0002 $\pm$ 0.0000	0.6976 $\pm$ 0.0046
MobileNetV3-Small	PlantVillage	cross-domain	3	2.7900 $\pm$ 0.0281	0.9296 $\pm$ 0.0068	0.1517 $\pm$ 0.0139	0.4799 $\pm$ 0.0438	0.6772 $\pm$ 0.0145	1.0000 $\pm$ 0.0000
MobileNetV3-Small	PlantVillage	cross-domain (T-scaled)	3	3.1303 $\pm$ 0.0525	1.0172 $\pm$ 0.0107	0.2938 $\pm$ 0.0128	0.5680 $\pm$ 0.0415	0.6769 $\pm$ 0.0145	0.6976 $\pm$ 0.0046
ResNet-50	PlantDoc	Core in-domain	3	1.2067 $\pm$ 0.0168	0.4769 $\pm$ 0.0172	0.1157 $\pm$ 0.0079	0.2413 $\pm$ 0.0669	0.1642 $\pm$ 0.0185	1.0000 $\pm$ 0.0000
ResNet-50	PlantDoc	Core in-domain (T-scaled)	3	1.1890 $\pm$ 0.0150	0.4699 $\pm$ 0.0187	0.0831 $\pm$ 0.0192	0.2254 $\pm$ 0.0373	0.1649 $\pm$ 0.0181	0.8780 $\pm$ 0.0447
EfficientNet-B0	PlantDoc	Core in-domain	3	1.1907 $\pm$ 0.0022	0.4858 $\pm$ 0.0065	0.1318 $\pm$ 0.0204	0.4597 $\pm$ 0.3660	0.1642 $\pm$ 0.0147	1.0000 $\pm$ 0.0000
EfficientNet-B0	PlantDoc	Core in-domain (T-scaled)	3	1.1332 $\pm$ 0.0075	0.4671 $\pm$ 0.0127	0.0659 $\pm$ 0.0183	0.2648 $\pm$ 0.0797	0.1632 $\pm$ 0.0129	0.7796 $\pm$ 0.0448
MobileNetV3-Small	PlantDoc	Core in-domain	3	1.5064 $\pm$ 0.0506	0.5952 $\pm$ 0.0185	0.1456 $\pm$ 0.0100	0.3238 $\pm$ 0.0464	0.2350 $\pm$ 0.0134	1.0000 $\pm$ 0.0000
MobileNetV3-Small	PlantDoc	Core in-domain (T-scaled)	3	1.4666 $\pm$ 0.0592	0.5822 $\pm$ 0.0215	0.0868 $\pm$ 0.0318	0.2306 $\pm$ 0.0163	0.2365 $\pm$ 0.0154	0.8317 $\pm$ 0.0512

Table 7. Backbone ranking by macro-F1 under each evaluation setting, per seed. A setting is STABLE when all seeds agree on the ordering. This is the direct test of whether the in-domain ranking predicts the cross-domain ranking, or whether an apparent reordering is within seed noise.

Setting	Seed Ranking (macro-F1; best first)	ResNet-50 macro-F1	EfficientNet-B0	macro-F1	MobileNet V3-Small macro-F1
PlantVillage in-domain	42 EfficientNet-B0 > MobileNet V3-Small > ResNet-50	0.9899	0.9956	0.9927	
PlantVillage in-domain	1337 EfficientNet-B0 > MobileNet V3-Small > ResNet-50	0.9925	0.9944	0.9931	
PlantVillage in-domain	2026 ResNet-50 > EfficientNet-B0 > MobileNet V3-Small	0.9957	0.9951	0.9923	
PlantVillage in-domain	all UNSTABLE: 2/3 seeds give EfficientNet-B0 > MobileNet V3-Small > ResNet-50	0.9927	0.9950	0.9927	
PlantVillage → PlantDoc Core	42 ResNet-50 > EfficientNet-B0 > MobileNet V3-Small	0.2414	0.2359	0.1982	
PlantVillage → PlantDoc Core	1337 EfficientNet-B0 > ResNet-50 > MobileNet V3-Small	0.2218	0.2312	0.1764	
PlantVillage → PlantDoc Core	2026 ResNet-50 > EfficientNet-B0 > MobileNet V3-Small	0.2242	0.2230	0.1862	
PlantVillage → PlantDoc Core	all UNSTABLE: 2/3 seeds give ResNet-50 > EfficientNet-B0 > MobileNet V3-Small	0.2291	0.2300	0.1869	
PlantDoc Core in-domain	42 ResNet-50 > EfficientNet-B0 > MobileNet V3-Small	0.6994	0.6927	0.5841	
PlantDoc Core in-domain	1337 ResNet-50 > EfficientNet-B0 > MobileNet V3-Small	0.7094	0.6663	0.5946	
PlantDoc Core in-domain	2026 EfficientNet-B0 > ResNet-50 > MobileNet V3-Small	0.6381	0.6716	0.5840	
PlantDoc Core in-domain	all UNSTABLE: 2/3 seeds give ResNet-50 > EfficientNet-B0 > MobileNet V3-Small	0.6823	0.6768	0.5876	

Table 8. Pairwise macro-F1 margins against seed noise. The margin is computed within each seed and then averaged. Pooled seed SD is the root-mean-square of the two models’ across-seed standard deviations. A margin comfortably exceeding the seed SD, with a consistent sign across all seeds, is a difference the seed alone does not explain; one below it is not claimed. With three seeds these are descriptive ratios, not significance tests.

Setting	Comparison	Seeds	Mean margin	Pooled seed SD	Margin / seed SD	Sign consistent
PlantVillage in-domain	ResNet-50 - EfficientNet-B0	3	-0.0023 $\pm$ 0.0032	0.0021	1.11	2/3
PlantVillage in-domain	ResNet-50 - MobileNetV3-Small	3	0.0000 $\pm$ 0.0031	0.0021	0.02	2/3
PlantVillage in-domain	EfficientNet-B0 - MobileNetV3-Small	3	0.0023 $\pm$ 0.0008	0.0005	4.73	3/3
PlantVillage $\rightarrow$ PlantDoc Core	ResNet-50 - EfficientNet-B0	3	-0.0009 $\pm$ 0.0077	0.0089	0.10	2/3
PlantVillage $\rightarrow$ PlantDoc Core	ResNet-50 - MobileNetV3-Small	3	0.0422 $\pm$ 0.0038	0.0108	3.90	3/3
PlantVillage $\rightarrow$ PlantDoc Core	EfficientNet-B0 - MobileNetV3-Small	3	0.0431 $\pm$ 0.0102	0.0090	4.78	3/3
PlantDoc Core in-domain	ResNet-50 - EfficientNet-B0	3	0.0054 $\pm$ 0.0383	0.0290	0.19	2/3
PlantDoc Core in-domain	ResNet-50 - MobileNetV3-Small	3	0.0947 $\pm$ 0.0352	0.0276	3.43	3/3
PlantDoc Core in-domain	EfficientNet-B0 - MobileNetV3-Small	3	0.0893 $\pm$ 0.0185	0.0108	8.30	3/3

Table 9. BCa bootstrap 95% confidence intervals for accuracy and macro-F1, seed 42, 10 000 resamples. Intervals describe sampling variability of the fixed evaluation set for one trained checkpoint; they are not seed variance, which is reported separately.

Setting	Model	N	Accuracy	Accuracy 95% CI	Macro-F1	Macro-F1 95% CI
PlantVillage in-domain	ResNet-50	10 709	0.9925	[0.9908, 0.9940]	0.9899	[0.9870, 0.9922]
PlantVillage in-domain	EfficientNet-B0	10 709	0.9968	[0.9956, 0.9978]	0.9956	[0.9938, 0.9969]
PlantVillage in-domain	MobileNetV3-Small	10 709	0.9950	[0.9935, 0.9962]	0.9927	[0.9901, 0.9948]
PlantDoc Core in-domain	ResNet-50	225	0.6933	[0.6311, 0.7511]	0.6994	[0.6526, 0.7721]
PlantDoc Core in-domain	EfficientNet-B0	225	0.6800	[0.6178, 0.7378]	0.6927	[0.6457, 0.7598]
PlantDoc Core in-domain	MobileNetV3-Small	225	0.5556	[0.4889, 0.6222]	0.5841	[0.5352, 0.6675]
PlantVillage $\rightarrow$ PlantDoc Core	ResNet-50	1 951	0.2778	[0.2584, 0.2978]	0.2414	[0.2253, 0.2595]
PlantVillage $\rightarrow$ PlantDoc Core	EfficientNet-B0	1 951	0.2742	[0.2553, 0.2942]	0.2359	[0.2203, 0.2538]
PlantVillage $\rightarrow$ PlantDoc Core	MobileNetV3-Small	1 951	0.2419	[0.2230, 0.2609]	0.1982	[0.1831, 0.2158]

Table 10. Pairwise McNemar exact tests on per-item correctness, seed 42. ‘A only’ and ‘B only’ are the discordant counts the test is computed from. Holm-adjusted p-values control the family-wise error rate across all pairwise tests reported here.

Setting	Comparison	Acc. A	Acc. B	Diff.	A only	B only	Discordant	p (exact)	p (Holm)
PlantVillage in-domain	ResNet-50 vs EfficientNet-B0	0.9925	0.9968	-0.0043	18	64	82	<0.0001	<0.0001
PlantVillage in-domain	ResNet-50 vs MobileNetV3-Small	0.9925	0.9950	-0.0024	34	60	94	0.0095	0.0363
PlantVillage in-domain	EfficientNet-B0 vs MobileNetV3-Small	0.9968	0.9950	+0.0019	37	17	54	0.0091	0.0363
PlantDoc Core in-domain	ResNet-50 vs EfficientNet-B0	0.6933	0.6800	+0.0133	20	17	37	0.7428	1.0000
PlantDoc Core in-domain	ResNet-50 vs MobileNetV3-Small	0.6933	0.5556	+0.1378	48	17	65	0.0002	0.0012
PlantDoc Core in-domain	EfficientNet-B0 vs MobileNetV3-Small	0.6800	0.5556	+0.1244	47	19	66	0.0008	0.0045
PlantVillage $\rightarrow$ PlantDoc Core	ResNet-50 vs EfficientNet-B0	0.2778	0.2742	+0.0036	179	172	351	0.7488	1.0000
PlantVillage $\rightarrow$ PlantDoc Core	ResNet-50 vs MobileNetV3-Small	0.2778	0.2419	+0.0359	207	137	344	0.0002	0.0013
PlantVillage $\rightarrow$ PlantDoc Core	EfficientNet-B0 vs MobileNetV3-Small	0.2742	0.2419	+0.0323	205	142	347	0.0008	0.0045

Table 11. Stability of each pairwise verdict across independent training seeds. Each seed’s McNemar test is computed on that seed’s own checkpoints and Holm-adjusted within that seed. A comparison significant in every seed is a conclusion about the architectures; one significant in a single seed is a statement about that run.

Setting	Comparison	Seeds	Significant (Holm)	Sign consistent	Mean acc. diff.
PlantVillage in-domain	ResNet-50 vs EfficientNet-B0	3	1/3	no	-0.0019
PlantVillage in-domain	ResNet-50 vs MobileNetV3-Small	3	2/3	no	-0.0003
PlantVillage in-domain	EfficientNet-B0 vs MobileNetV3-Small	3	2/3	yes	+0.0016
PlantDoc Core in-domain	ResNet-50 vs EfficientNet-B0	3	0/3	no	-0.0015
PlantDoc Core in-domain	ResNet-50 vs MobileNetV3-Small	3	2/3	yes	+0.1037
PlantDoc Core in-domain	EfficientNet-B0 vs MobileNetV3-Small	3	3/3	yes	+0.1052
PlantVillage $\rightarrow$ PlantDoc Core	ResNet-50 vs EfficientNet-B0	3	0/3	no	-0.0070
PlantVillage $\rightarrow$ PlantDoc Core	ResNet-50 vs MobileNetV3-Small	3	3/3	yes	+0.0345
PlantVillage $\rightarrow$ PlantDoc Core	EfficientNet-B0 vs MobileNetV3-Small	3	3/3	yes	+0.0415

8 Risk–Coverage and Seed Spread




Fig. 2. Confidence–abstention curves, seed-averaged on a common coverage grid with a $\pm$ one standard deviation band.




Fig. 3. Backbone macro-F1 by evaluation setting: mean with a seed standard deviation bar, and each individual seed drawn beside it. With three seeds the individual points carry more information than the interval computed from them.

9 Per-Run Diagnostics

Confusion matrices and reliability diagrams are generated for every run and evaluation. They are released as individual files rather than reproduced here, since the full set runs to dozens of figures:

- experiments/ica26/figures/fig_confusion_<run>_<eval>.pdf
- experiments/ica26/figures/fig_reliability_<run>.pdf

10 Per-Seed Values

Every aggregated table above has a per-seed companion in CSV, so each reported mean and standard deviation can be recomputed without rerunning any experiment:

- table2_in_domain_performance_per_seed.csv
- table3_cross_domain_performance_per_seed.csv
- table4_efficiency_per_seed.csv
- table5_domain_shift_degradation_per_seed.csv
- table6_calibration_per_seed.csv